

Parallel Machine Learning and Artificial Intelligence

Dr. Handan Liu

h.liu@northeastern.edu

Northeastern University

Content

- High Performance Parallel Computing


- Overview (done)
- Concepts and Terminology (done)
- Parallel Computer Memory Architectures (done)
- Parallel Programming Model (done)
 - ✓ Parallel Implementations – OpenMP Programming
 - ✓ Parallel Implementations – MPI Programming
- Parallel Examples and Exercises

OpenMP Directives

- Work-Sharing Constructs

- A work-sharing construct divides the execution of the enclosed code region among the members of the team that encounter it.
- Work-sharing constructs do not launch new threads
- There is no implied barrier upon entry to a work-sharing construct, however there is an implied barrier at the end of a work sharing construct.

Types of Work-Sharing Constructs:



DO / for - shares iterations of a loop across the team.
Represents a type of "data parallelism".

SECTIONS - breaks work into separate, discrete sections. Each section is executed by a thread. Can be used to implement a type of "functional parallelism".

SINGLE - serializes a section of code.

Restrictions:

- A work-sharing construct must be enclosed dynamically within a parallel region in order for the directive to execute in parallel.
- Work-sharing constructs must be encountered by all members of a team or none at all
- Successive work-sharing constructs must be encountered in the same order by all members of a team

DO / for Directive

- Purpose:

- The **DO / for** directive specifies that the iterations of the loop immediately following it must be executed in parallel by the team. This assumes a parallel region has already been initiated, otherwise it executes in serial on a single processor.

- Format:

C/C++

#pragma omp for *[clause ...]* newline

schedule (type [,chunk])

ordered

private (*list*)

firstprivate (*list*)

lastprivate (*list*)

shared (*list*)

reduction (*operator: list*)

collapse (*n*)

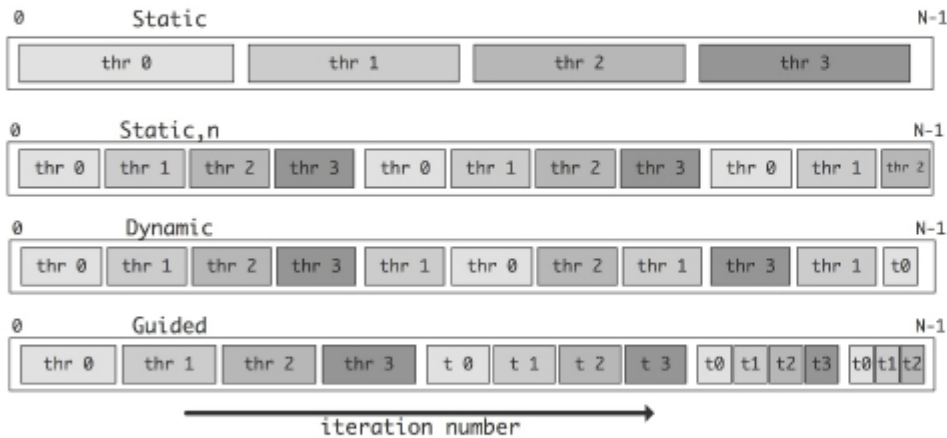
nowait

Important

for_loop

OpenMP: For Loop & Scheduling

`#pragma omp parallel for schedule(scheduling-type)`



- `schedule(static, chunk-size)`
- `schedule(dynamic, chunk-size)`
- `schedule(guided, chunk-size)`
- `schedule(auto)`

Example: DO / for Directive – Vector Addition

```
// C / C++ - for Directive Example
```

```
#define N 1000
```

```
main(int argc, char *argv[]) {
```

```
    int i;  
    float a[N], b[N], c[N];
```

```
    /* Some initializations */
```

```
    for (i=0; i < N; i++)  
        a[i] = b[i] = i * 1.0;
```

```
    {  
  
        for (i=0; i < N; i++)  
            c[i] = a[i] + b[i];  
    }  
}
```

```
// C / C++ - for Directive Example
```

```
#include <omp.h>
```

```
#define N 1000
```

```
#define CHUNKSIZE 100
```

```
main(int argc, char *argv[]) {
```

```
    int i, chunk;  
    float a[N], b[N], c[N];
```

```
    /* Some initializations */
```

```
    for (i=0; i < N; i++)  
        a[i] = b[i] = i * 1.0;  
    chunk = CHUNKSIZE;
```

```
    #pragma omp parallel shared(a,b,c,chunk,N) private(i)
```

```
    {  
  
        #pragma omp for schedule(dynamic,chunk) nowait  
        for (i=0; i < N; i++)  
            c[i] = a[i] + b[i];  
    } /* end of parallel region */  
}
```

OpenMP Directives

- Data Scope Attribute Clauses
 - Global (shared) variables:
 - ✓ File scope variables and static (in C)
 - Private variables:
 - ✓ Loop index variables
 - ✓ Stack variables in subroutines called from parallel regions

<https://www.openmp.org/spec-html/5.1/openmps6.html#x13-120001.2.6>

Data Scope Attribute Clauses

- The OpenMP Data Scope Attribute Clauses are used to explicitly define how variables should be scoped. They include:
 - PRIVATE
 - FIRSTPRIVATE
 - LASTPRIVATE
 - SHARED
 - DEFAULT
 - REDUCTION
 - COPYIN

PRIVATE Clause

- Purpose:
 - The PRIVATE clause declares variables in its list to be private to each thread.
- PRIVATE variables behave as follows:
 - A new object of the same type is declared once for each thread in the team
 - All references to the original object are replaced with references to the new object
 - Should be assumed to be uninitialized for each thread

SHARED Clause

- Purpose:

- The SHARED clause declares variables in its list to be shared among all threads in the team.

- Notes:

- A shared variable exists in only one memory location and all threads can read or write to that address
- It is the programmer's responsibility to ensure that multiple threads properly access SHARED variables

Data Scope Attribute Clauses

- Data Scope Attribute Clauses are used in conjunction with several directives (**PARALLEL**, **DO/for**, and **SECTIONS**) to control the scoping of enclosed variables.
- These constructs provide the ability to control the data environment during execution of parallel constructs.
- Data Scope Attribute Clauses are effective only within their lexical/static extent.

Example: Shared / Private Variables

```
// C / C++ - for Directive Example
#include <omp.h>
#define N 1000
#define CHUNKSIZE 100

main(int argc, char *argv[]) {
    int i, chunk;
    float a[N], b[N], c[N];

    /* Some initializations */
    for (i=0; i < N; i++)
        a[i] = b[i] = i * 1.0;
    chunk = CHUNKSIZE;

    #pragma omp parallel shared(a,b,c,chunk,N) private(i)
    {
        #pragma omp for schedule(dynamic,chunk) nowait
        for (i=0; i < N; i++)
            c[i] = a[i] + b[i];
    } /* end of parallel region */
}
```

Example: Shared / Private Variables

```
#pragma omp parallel for  
{  
    for(i=1; i<=n; i++){  
        temp = 2.0*a[i];  
        a[i] = temp;  
        b[i] = c[i]/temp;  
    }  
}
```

Loop level Parallelization

- Requirements for Loop Parallelization

- no dependencies between loop indices
- an element of an array is assigned to by at most one iteration
- no loop iteration reads array elements modified by any other dependency
- due to overhead of parallelization - use only on loops where individual iterations take a long time

```
#pragma omp parallel for
for(i=1; i<=n; i++)
    a[i] = b[i] + c[i]
```

```
#pragma omp parallel for
for(i=2; i<=5; i++)
    a[i] = a[i] + a[i-1];
```

Quiz 1

- Will take a quiz on Feb 5th, Friday.
- Start at 11AM.
- Online on Canvas -> Quizzes
- Most are choice questions or multiple-choice questions.
- Everyone MUST turn on the camera; otherwise the quiz score will be treated as 0.

- Stay safe!
- See you next class!

Next Lecture will Continue:

Hands-on Lab: OpenMP Programming



